


Parallel Computing with MATLAB

Jiro Doke, Ph.D.
Senior Application Engineer

Sarah Wait Zaranek, Ph.D.
MATLAB Product Marketing




A Question to Consider

Do you want to speed up your algorithms or deal with large data?

If so...


- Do you have a multi-core or multi-processor computer?
- Do you have a high-end graphics processing unit (GPU)?
- Do you have access to a computer cluster?



Utilizing Additional Processing Power

- Built-in multithreading (implicit)
 - Core MATLAB and Image Processing Toolbox
 - Utility for specific matrix operations (linear algebra, fft, filter, etc)
 - No necessary code change
- Parallel computing tools (explicit)
 - Parallel Computing Toolbox
 - MATLAB Distributed Computing Server
 - Broad utility controlled by the MATLAB user

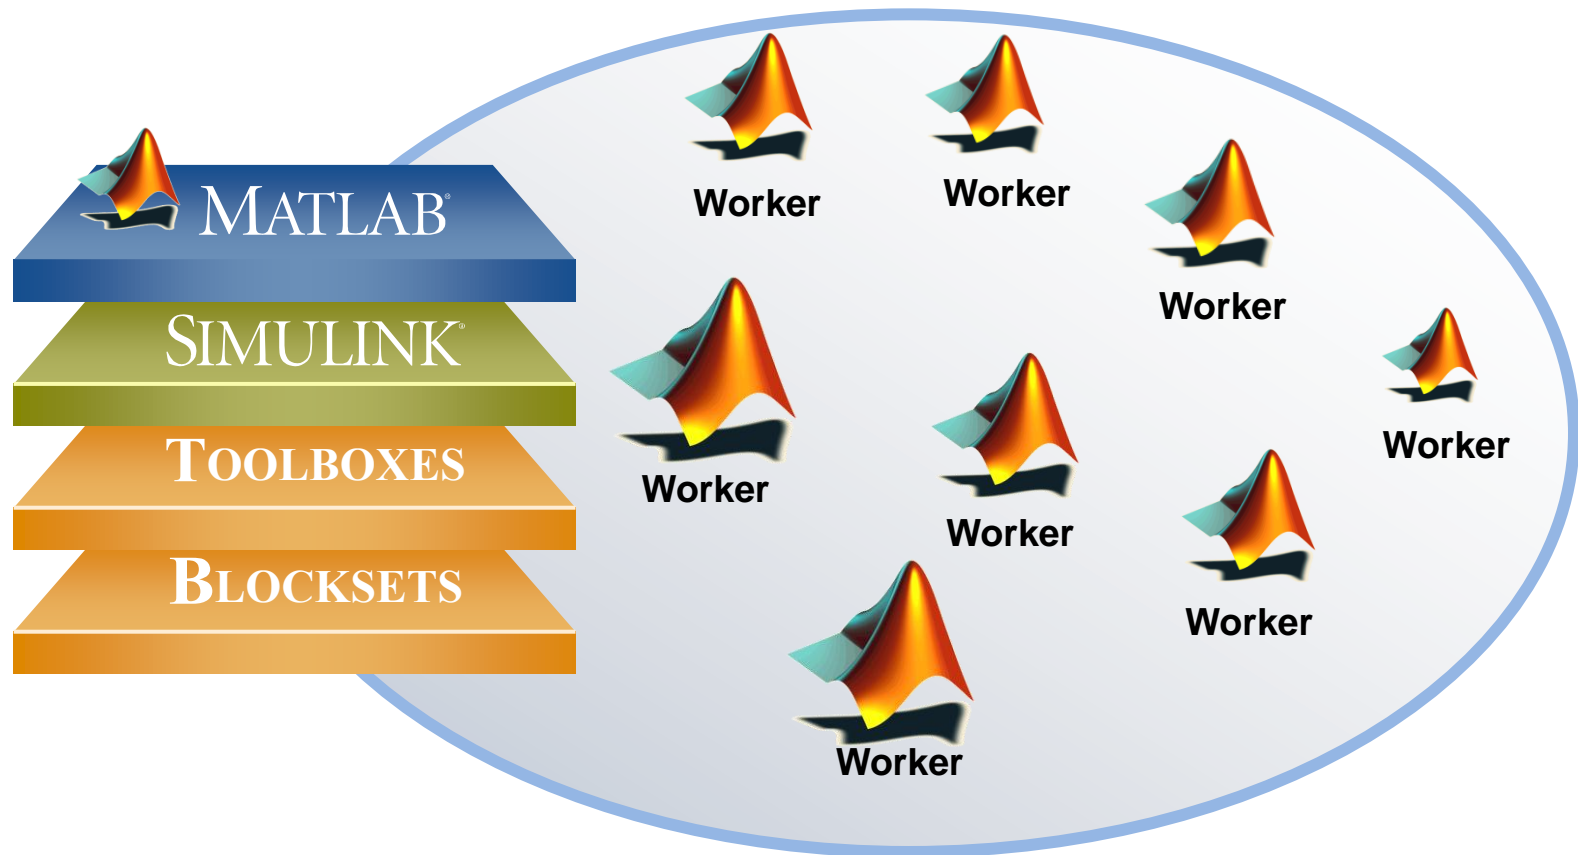
Agenda




Introduction to Parallel Computing Tools

- Using Multi-core/Multi-processor Machines
- Using Graphics Processing Units (GPUs)
- Scaling Up to a Cluster

Going Beyond Serial MATLAB Applications

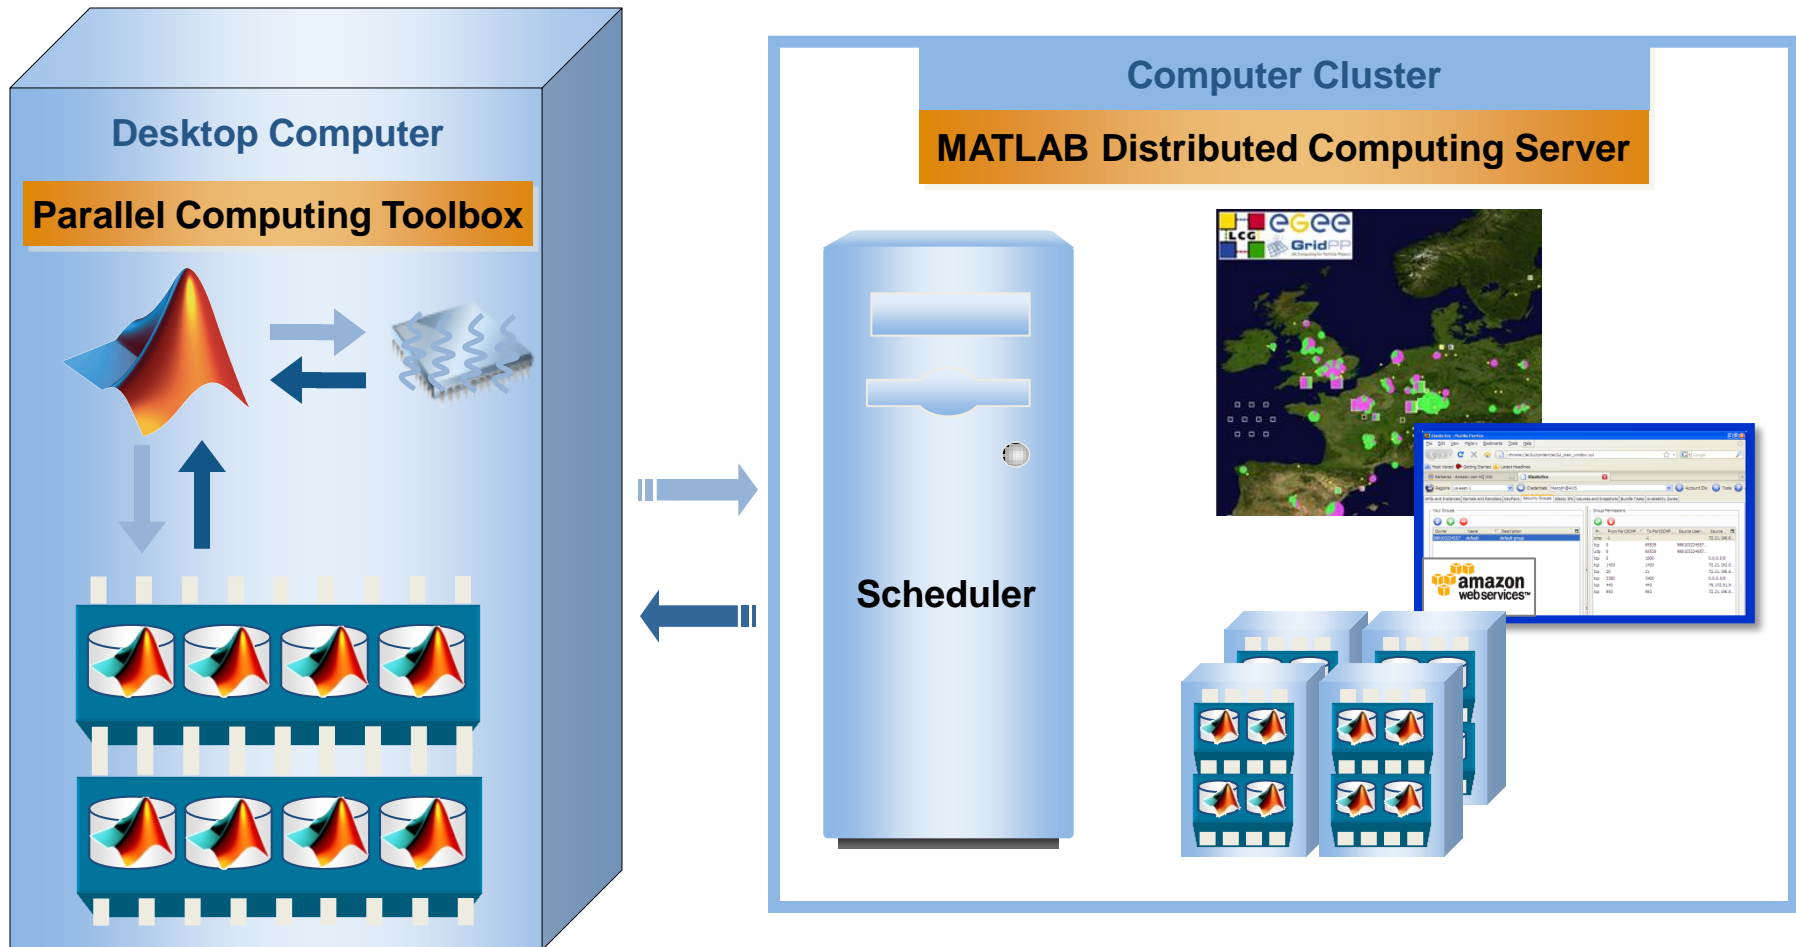


Parallel Computing on the Desktop



- Use Parallel Computing Toolbox
- Speed up parallel applications on local computer
- Take full advantage of desktop power by using CPUs and GPUs (up to 12 workers in R2011b)
- Separate computer cluster not required


Scale Up to Clusters, Grids and Clouds



Parallel Computing enables you to ...


Larger Compute Pool

Speed up Computations




Larger Memory Pool

Work with Large Data



Agenda

- Introduction to Parallel Computing Tools



Using Multi-core/Multi-processor Machines

- Using Graphics Processing Units (GPUs)
- Scaling Up to a Cluster

Programming Parallel Applications



Using Additional Cores/Processors (CPUs)



Ease of Use

- Support built into Toolboxes





Greater Control

Example:


Built-in Support for Parallelism in Other Tools

- Use built-in support for Parallel Computing Toolbox in Optimization Toolbox
- Run optimization in parallel
- Use pool of MATLAB workers



Other Tools Providing Parallel Computing Support

- Optimization Toolbox
- Global Optimization Toolbox
- Statistics Toolbox
- Simulink Design Optimization
- Bioinformatics Toolbox
- Model-Based Calibration Toolbox
- ...



<http://www.mathworks.com/products/parallel-computing/builtin-parallel-support.html>

Using Additional Cores/Processors (CPUs)



Ease of Use



- Support built into Toolboxes
- Simple programming constructs:
`parfor`



Greater Control

Running Independent Tasks or Iterations


- Ideal problem for parallel computing
- No dependencies or communications between tasks
- Examples include parameter sweeps and Monte Carlo simulations



Example:

Parameter Sweep of ODEs

- Parameter sweep of ODE system
- Use pool of MATLAB workers
- Convert `for` to `parfor`
- Interleave serial and parallel code



Damped spring oscillator

$$\underbrace{m}_{5} \ddot{x} + \underbrace{b}_{1,2,\dots} \dot{x} + \underbrace{k}_{1,2,\dots} x = 0$$

- Sweep through different values of **b** and **k**
- Record peak value for each simulation

Tips for using `parfor`

- Requirement: Task and order independence
- Classification of Variables
 - One of the most common type of problems people run into when working with PARFOR.
 - At runtime, MATLAB needs determine how each variable would get treated.
 - Documentation: Parallel Computing Toolbox → User's Guide → Parallel for-Loops → Advanced Topics
- <http://blogs.mathworks.com/loren/2009/10/02/using-parfor-loops-getting-up-and-running/>

Using Additional Cores/Processors (CPUs)



Ease of Use


- Support built into Toolboxes
- Simple programming constructs:
`parfor`
- Full control of parallelization:
jobs and tasks



Greater Control

Agenda

- Introduction to Parallel Computing Tools
- Using Multi-core/Multi-processor Machines




Using Graphics Processing Units (GPUs)


- Scaling Up to a Cluster

Gaining Performance with More Hardware

Using More Cores (CPUs)



Using GPUs



What is a Graphics Processing Unit (GPU)

- Originally for graphics acceleration, now also used for scientific calculations
- Massively parallel array of integer and floating point processors
 - Typically hundreds of processors per card
 - GPU cores complement CPU cores
- Dedicated high-speed memory




* Parallel Computing Toolbox requires NVIDIA GPUs with Compute Capability 1.3 or greater, including NVIDIA Tesla 10-series and 20-series products. See http://www.nvidia.com/object/cuda_gpus.html for a complete listing

Example:

GPU Computing in the Parallel Computing Toolbox


- Solve 2nd order wave equation:
$$\frac{\partial^2 u}{\partial t^2} = \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2}$$
- Send and create data on the GPU
- Run calculations with built-in GPU functions

Solution of 2nd Order Wave Equation



Benchmark: Solving 2D Wave Equation

CPU vs GPU



Summary of Options for Targeting GPUs



Ease of Use

- Use GPU array interface with MATLAB built-in functions
- Execute custom functions on elements of the GPU array
- Create kernels from existing CUDA code and PTX files




Greater Control

Webinar: “GPU Computing with MATLAB”

<http://www.mathworks.com/company/events/webinars/wbnr59816.html>

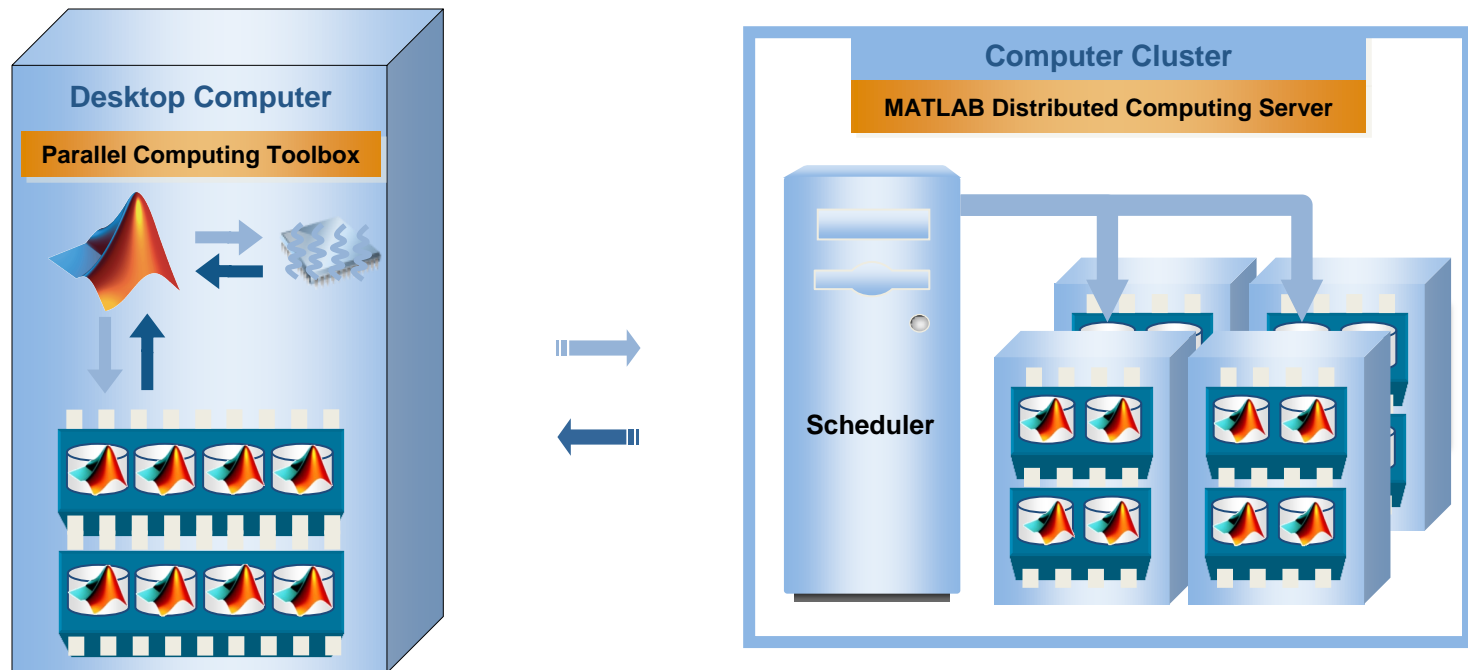
Agenda

- Introduction to Parallel Computing Tools
- Using Multi-core/Multi-processor Machines
- Using Graphics Processing Units (GPUs)



Scaling Up to a Cluster

Setting Up Cluster Computing (for System Admins)



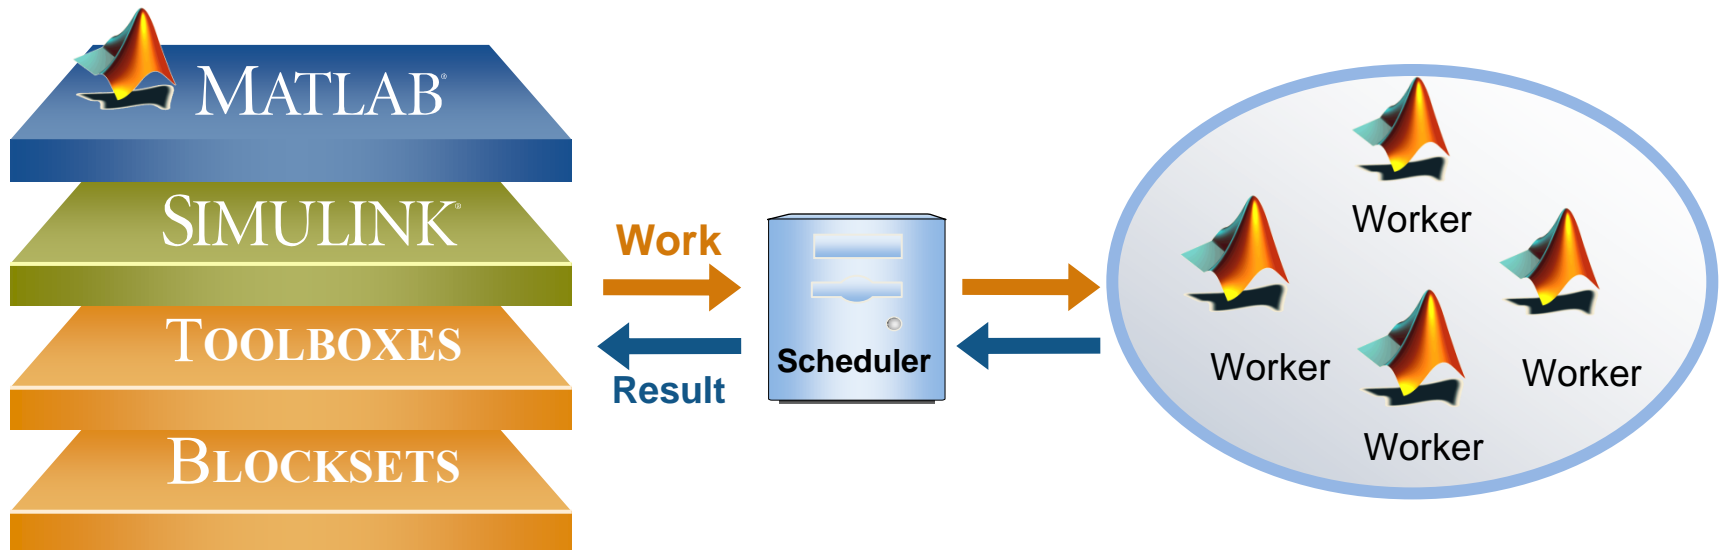
MATLAB Distributed Computing Server

- All-product install
- Worker license per process
- License by packs: 8, 16, 32, 64, etc.
- No additional toolbox licenses needed

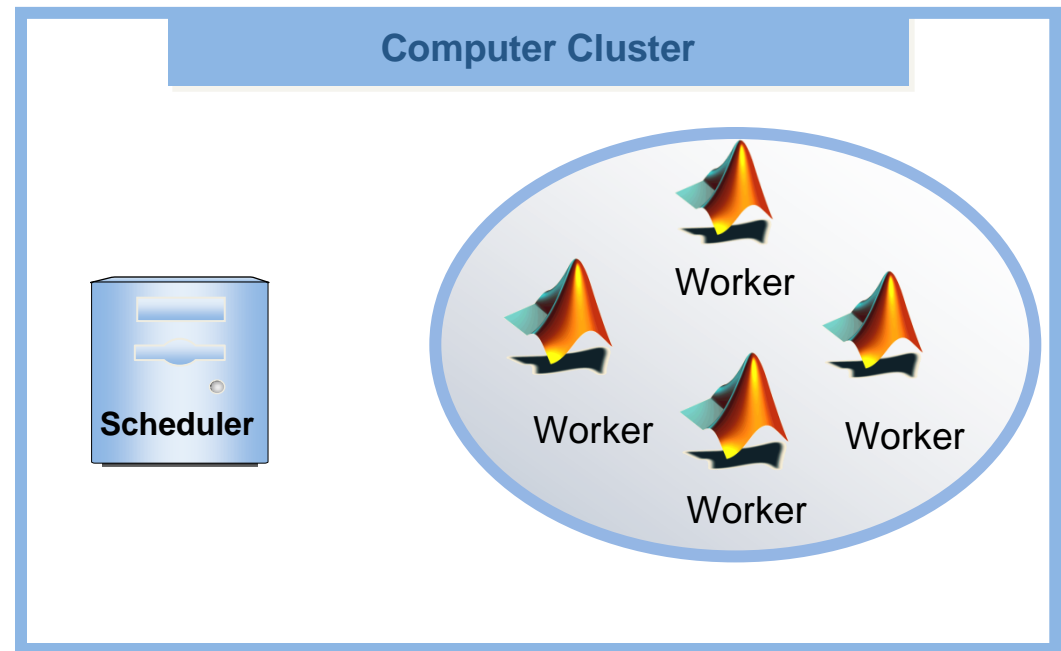
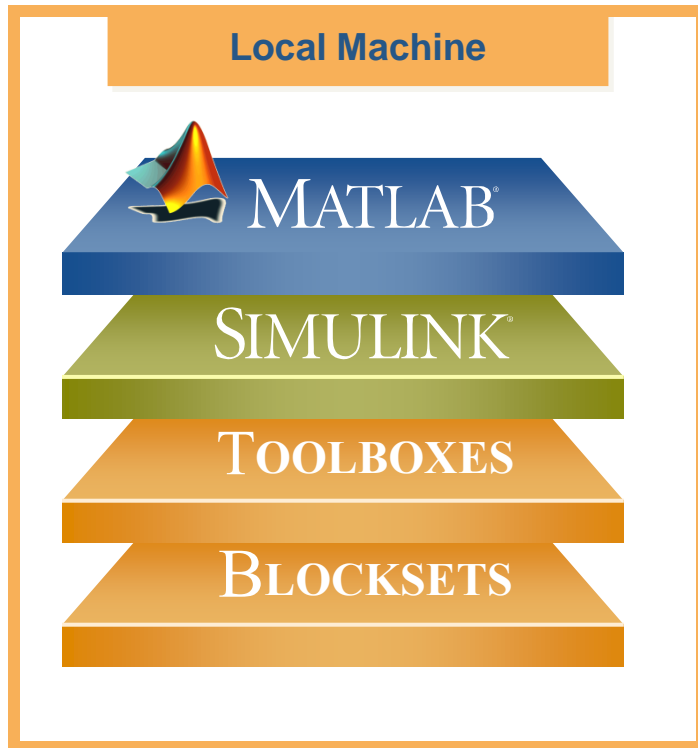
Why scale up to a cluster?

- Solve larger, computationally-intensive problems with more processing power
- Solve memory-intensive problems
- Schedule computations to offload from your local machine

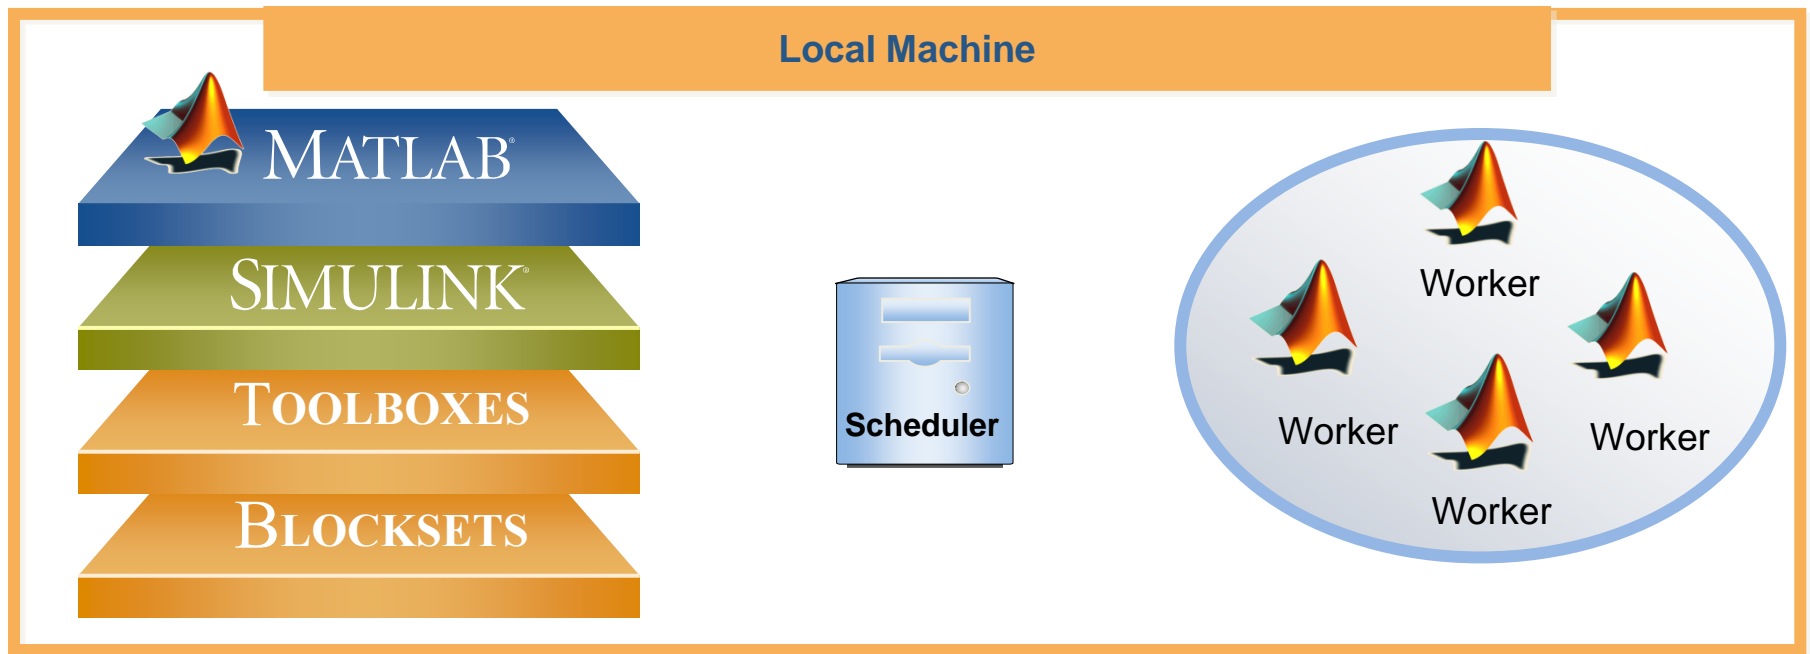
Scheduling Work (batch)



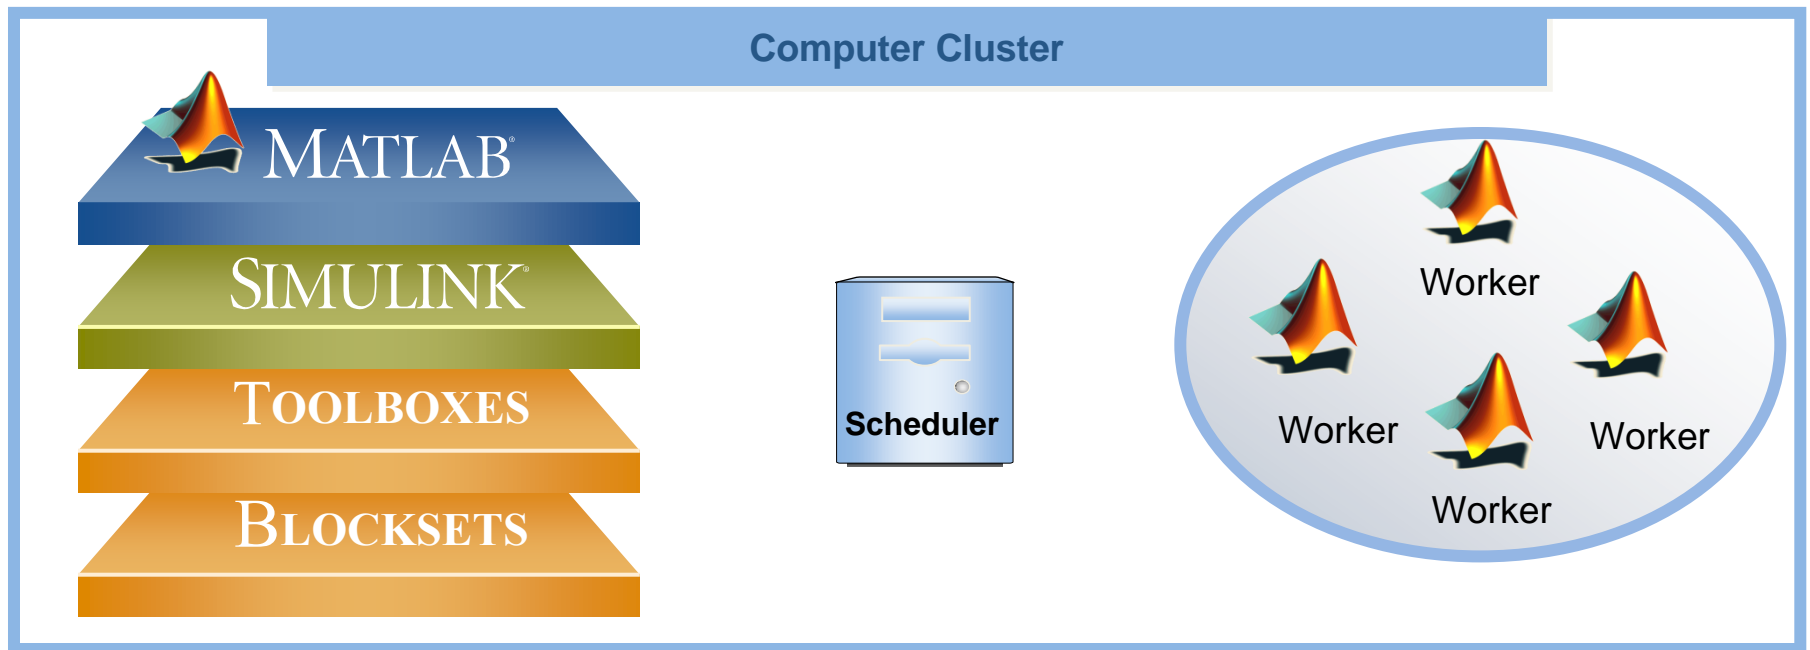
Scheduling Work (batch)



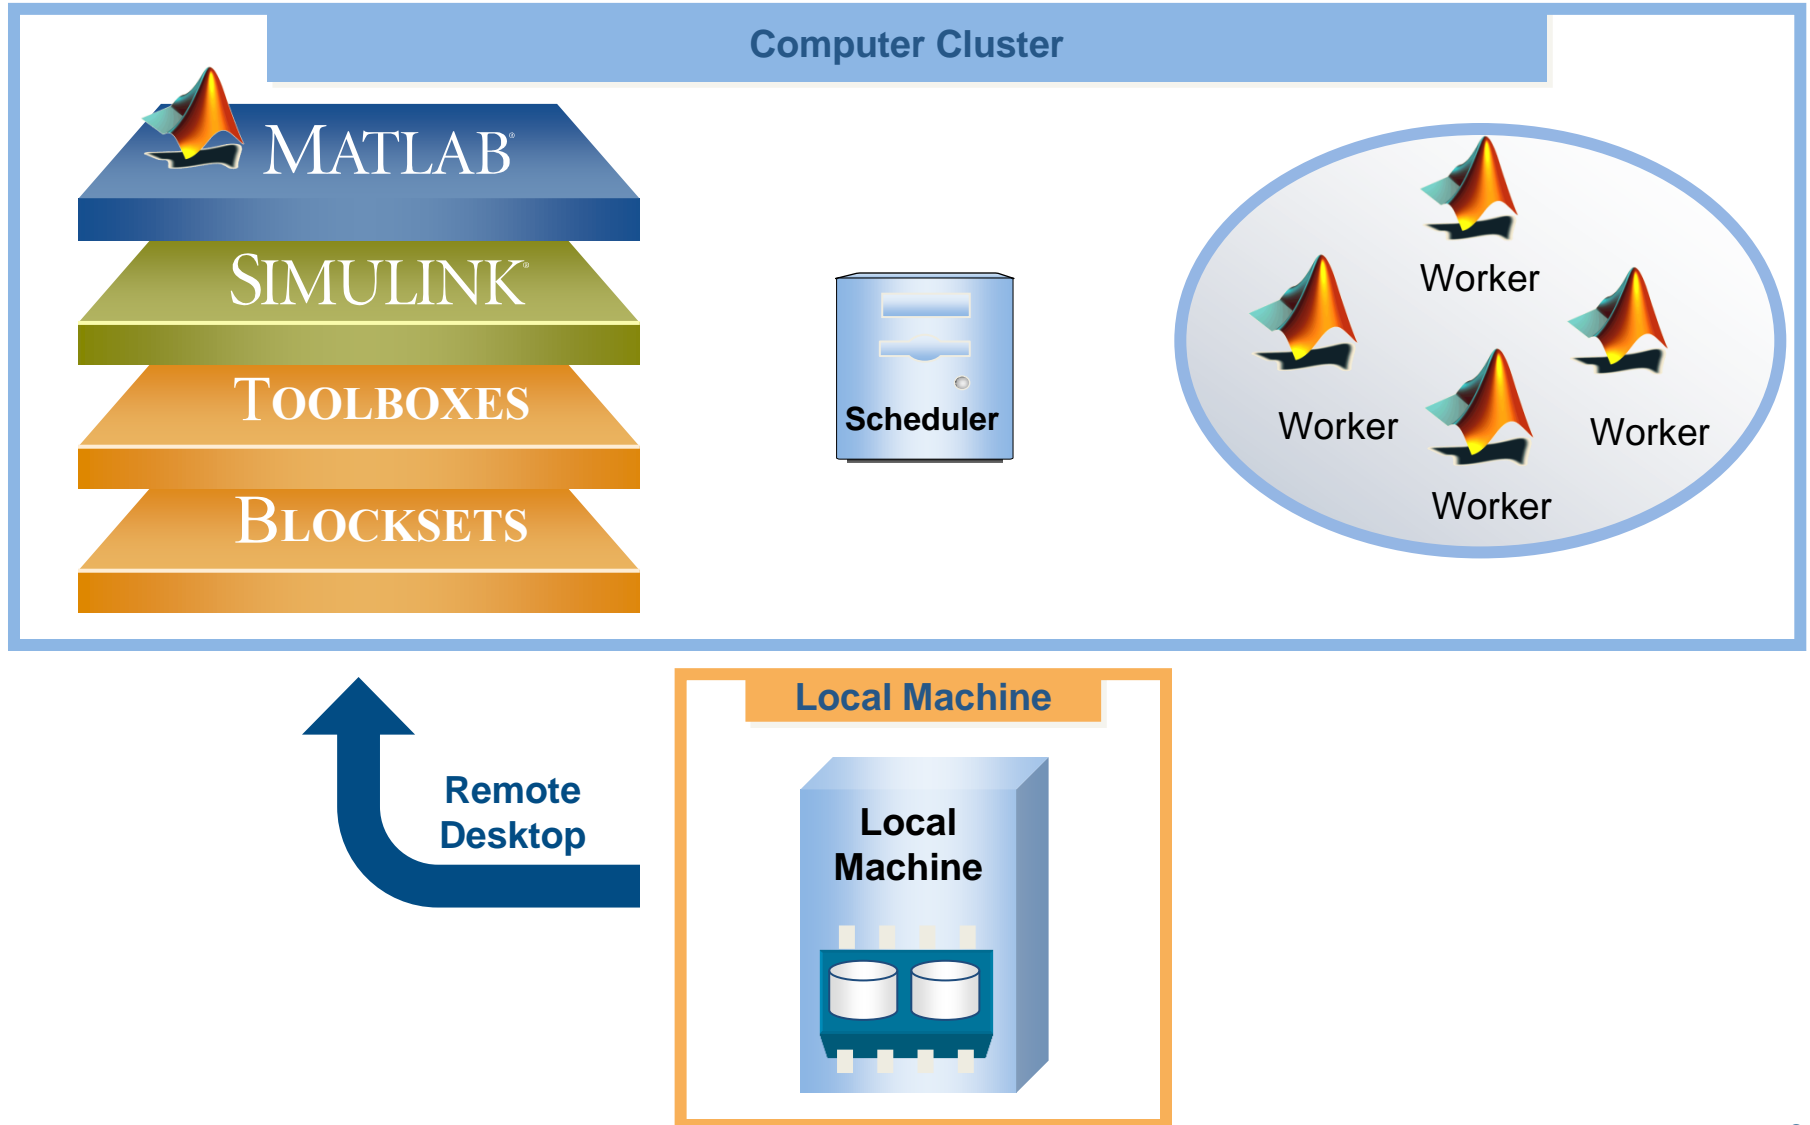
Scheduling Work (batch)



Scheduling Work (batch)




Scheduling Work (batch)



Example: Scheduled Processing

- Offload processing to workers (local or cluster)
- Regain control of MATLAB after offloading
- Monitor progress of scheduled job
- Retrieve results from job

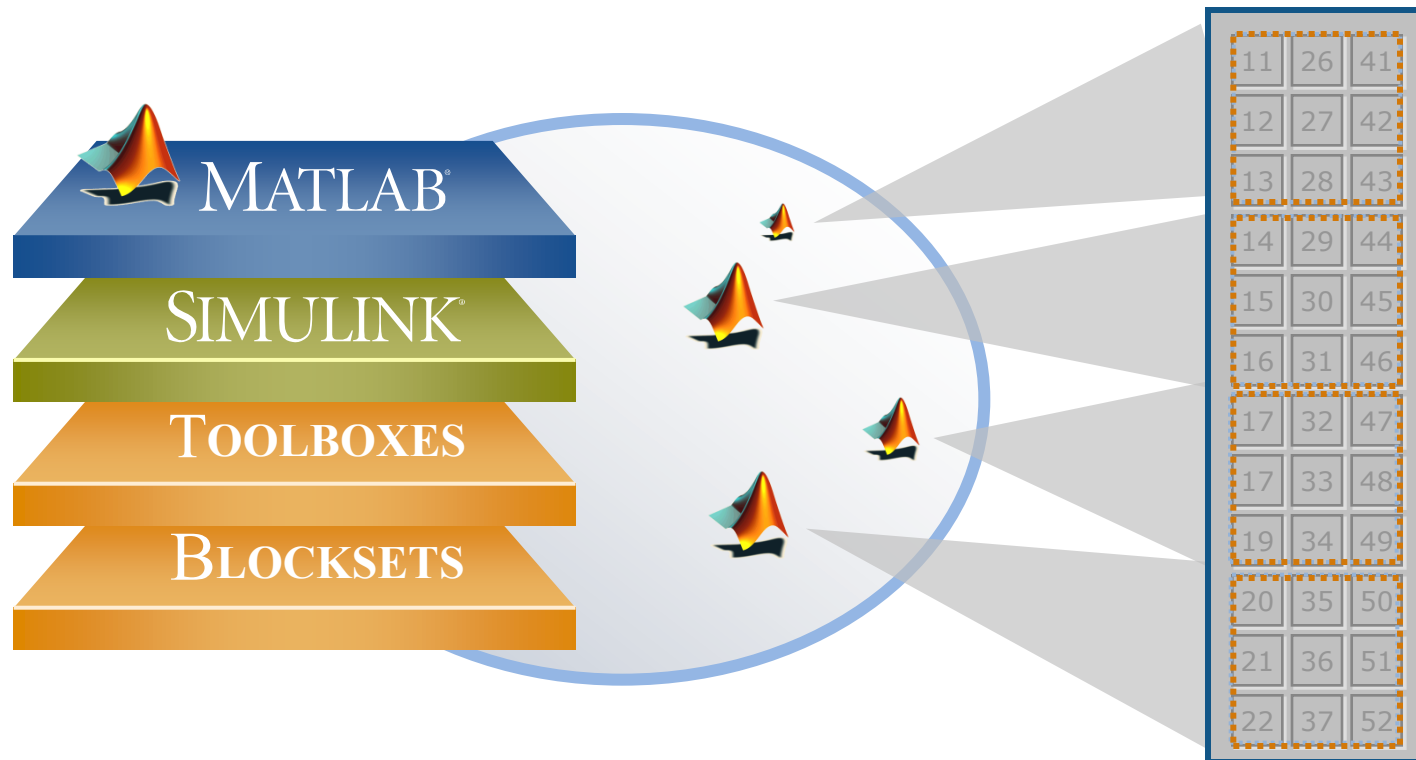


Damped spring oscillator

$$\underbrace{5}_{m} \ddot{x} + \underbrace{b}_{1,2,\dots} \dot{x} + \underbrace{k}_{1,2,\dots} x = 0$$

- Sweep through different values of **b** and **k**
- Record peak value for each simulation

Distributing Large Data



Remotely Manipulate Array
from Client MATLAB

Distributed Array
Lives on the Workers

Client-side Distributed Arrays and SPMD

- Client-side distributed arrays
 - Class **distributed**
 - Can be created and manipulated directly from the client.
 - Simpler access to memory on labs
 - 100s of built-in functions that operate on distributed arrays, including client-side visualization functions

- **spmd**
 - Block of code executed on workers
 - Worker specific commands
 - Explicit communication between workers using MPI
 - Mixture of parallel and serial code

Summary

- Speed up parallel applications on desktop by using *Parallel Computing Toolbox*
- Take full advantage of CPU and GPU hardware
- Use *MATLAB Distributed Computing Server* to
 - Scale up to clusters, grids and clouds
 - Work with data that is too large to fit on to desktop computers

For more information

Visit

www.mathworks.com/products/parallel-computing